



Classification of Diabetic Retinopathy using Deep Learning


Nick Weda
COMP 542 : Machine Learning



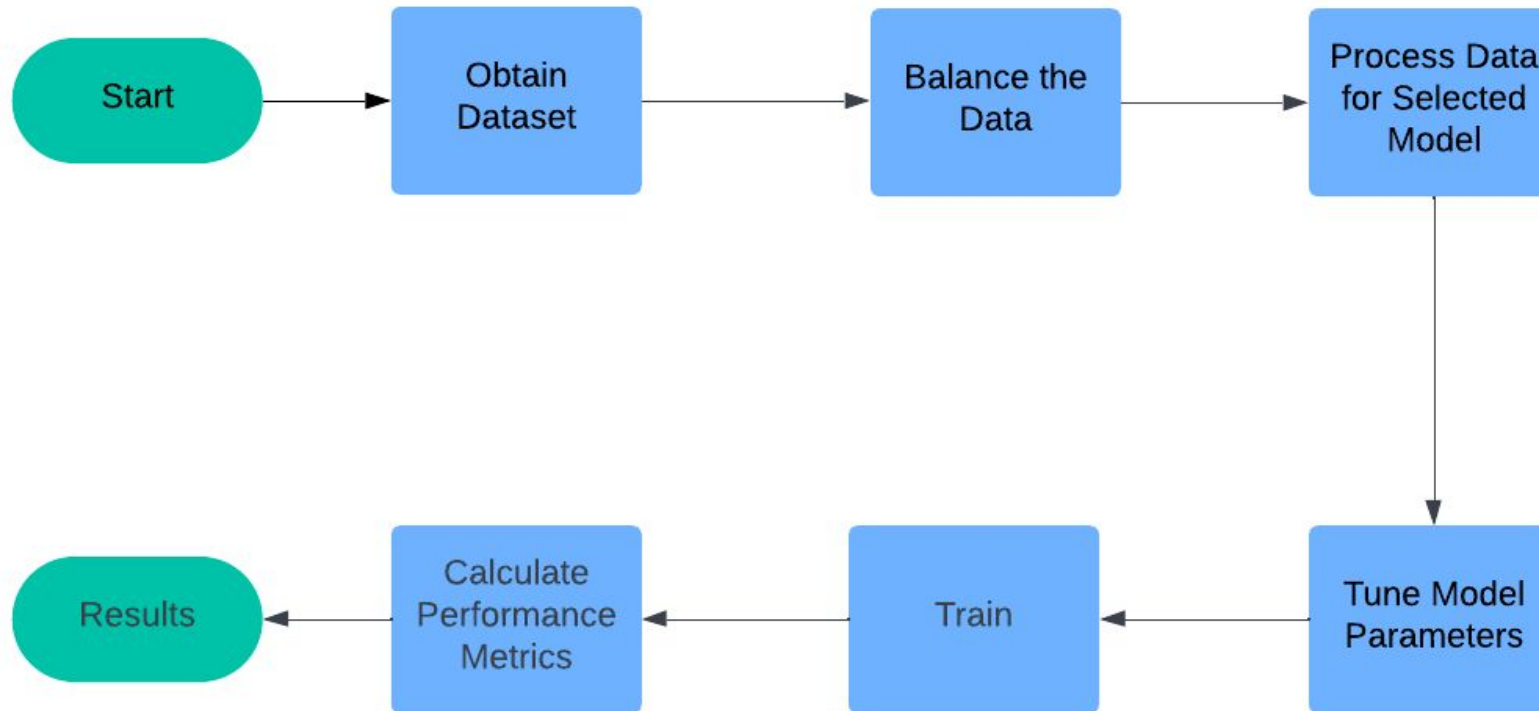
Department of Computer Science
California State University, Northridge

INTRODUCTION

- Diabetic Retinopathy is a complication of diabetes that can lead to blindness
- The disease is manageable if discovered and treated early
- Machine Learning can assist in the detection of early cases



METHOD



DATASET

- Dataset source
- Kaggle Competition - 2015 - \$100,000 prize
- EyePacs - Provider of Data - Diabetic Retinopathy Screening Program in California



Department of Computer Science
California State University, Northridge

DATASET

- Dataset Size
- 37.9 GB
- 35,126 images.
- 17,563 people. Left and Right eye.
- High resolution images of varying size. 4k+
- Accessing the Dataset:

Wrote a small script utilizing command line tools to save to the cloud



Recommended

Basic (100 GB)

~~\$1.99~~ \$0.49/mo

for 2 months

Save up to \$3 with offer

Get discount

\$1.99/month after

✓ 100 GB of storage for
Photos, Drive & Gmail

✓ Share storage with up to 5
others






DATASET

- Data Exploration
- The disease is on a severity scale of 0-4, where 0 is healthy and 4 is severe

File: 16_left | Label: 4




File: 10_left | Label: 0




File: 15_left | Label: 1



File: 15_right | Label: 2




File: 99_left | Label: 3



DATASET


- Number of Samples
- Data was hand labeled by medical professionals
- Dataset is highly imbalanced
- Characteristic of Medical Data
- Most surveyed patients are healthy

matplotlib



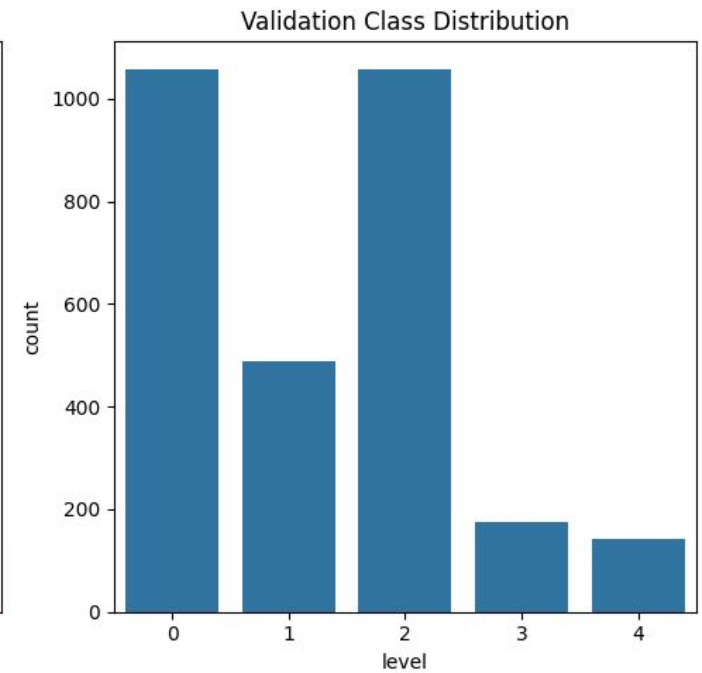
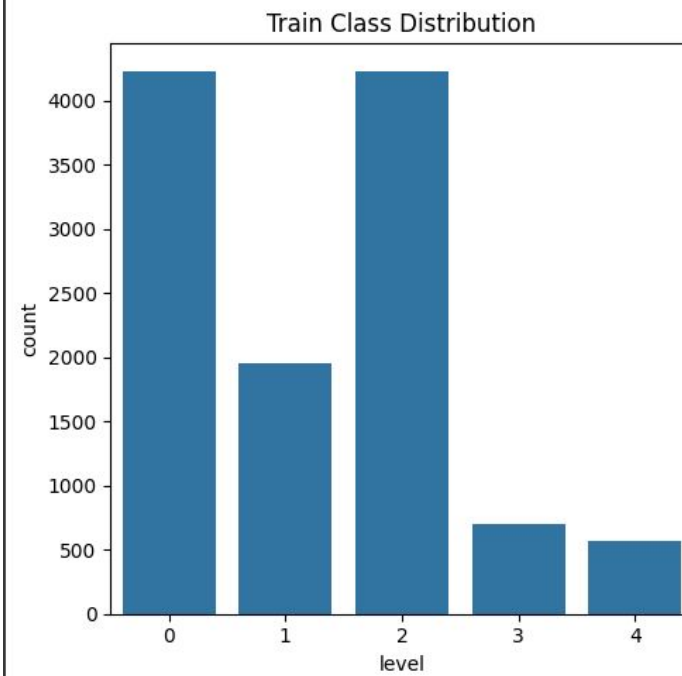
DATASET

- Combatting Imbalance
- Undersample the majority class
- Use the Scikit Imbalanced-Learn Library
- Randomly select majority class examples to keep
- New Dataset Size: 14,608 (~40% Decrease)



DATASET

- Combatting Imbalance
- Train-Test Split
- Stratified Sampling to maintain the data's distribution



Choice of Model



- Keras vs. PyTorch APIs
- Keras has a simpler interface
- Includes a set of Convolutional Neural Network architectures
- Based on the dataset size:
- ResNet-50, MobileNetV2, **EfficientNetB0 - 237 Layers**




Image Pre-Processing

- All 3 previously mentioned models require 224x224 images
- Careful cropping to avoid distorting the image
- Crop to **Region Of Interest** (Capture the retina)
- Retain **Aspect Ratio** (Width/Height Proportion)
- Add Black Border if necessary to preserve aspect ratio





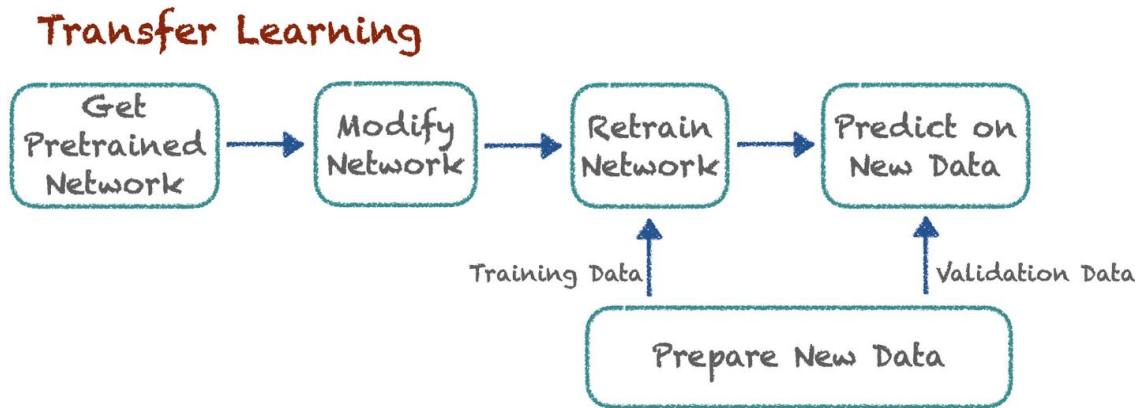
Image Pre-Processing

- We can take the opportunity to apply some light **augmentations**
- Random Horizontal Flip, Random Brightness, Random Zoom
- Save future computation by transforming once and then saving



Model Parameters


- Use pre-trained weights on ImageNet. Use **Transfer Learning**. Saves computation.
- Freeze beginning layers, basic pattern recognition.
- Weights will not be updated.
- Unfreeze the last layers (50/237), domain specific application.



Model Parameters

- Transfer learning adjusts what our **learning rate** should be. How much the weights change on each pass.
- Conduct a Learning Rate Range Test (LRRT). Tests the model on several rates.
- Ideal in this case = 0.0001 (1e-4)
- **Loss** is the difference between predicted and actual and should decrease over time


KT Keras Tuner



Department of Computer Science
California State University, Northridge


Model Parameters

- Other tuning parameters
- **Batch Size** - How many images fed in at once. Higher = faster training. (128)
- **Epochs** - How many entire passes of the training data through the network. (15)



Control Overfitting

- EfficientNetB0 is a fairly complex model. 14,000 images on on the smaller end for it.
- **Dropout:** Randomly dropout a fraction of neurons on every forward pass (20%)
- **L2 Regularization:** Add a penalty term to the loss function (cross entropy), pushing the weights to a smaller value



L2 Regularization

$$\text{Cost} = \underbrace{\sum_{i=0}^N (y_i - \sum_{j=0}^M x_{ij} W_j)^2}_{\text{Loss function}} + \lambda \underbrace{\sum_{j=0}^M W_j^2}_{\text{Regularization Term}}$$



RESULT


- Performance Metrics
- Accuracy = $(TP + TN) / \text{Total}$
- Precision (of believed positive) = $TP / (TP + FP)$
- Recall (of actual positive) = $TP / (TP + FN)$
- F1 Score - Harmonic Mean
- Support - occurrence of class
- Macro Average - across each class
- Weighted Average - across each class (proportion)

Classification Report:					
	precision	recall	f1-score	support	
0	0.5504	0.6248	0.5852	1058	
1	0.2767	0.2331	0.2531	489	
2	0.5504	0.5520	0.5512	1058	
3	0.4500	0.3600	0.4000	175	
4	0.6019	0.4577	0.5200	142	
accuracy			0.5089	2922	
macro avg	0.4859	0.4455	0.4619	2922	
weighted avg	0.5011	0.5089	0.5031	2922	




RESULT

- Performance Metrics
- Only predicted the dominating classes with the most accuracy.
- Need more data for 1,3,4.
- It doesn't exist for 3,4.
- Training time ~30 min.
- Model was overfitting, train accuracy was high at 87%+.



RESULT

- Possible Improvements:
- Build a network from scratch instead of relying on transfer learning.
- Implement proper SMOTE oversampling of images, creating synthetic images.
- Apply augmentations on the fly for more variability of data.
- Develop a novel feature extraction algorithm that can be used for accurate classification.



Closing Remarks



Resources X

You are subscribed to Colab Pro. [Learn more](#)

Available: 0.52 compute units

Usage rate: approximately 0 per hour

You have 0 active sessions.

[Manage sessions](#)



Department of Computer Science
California State University, Northridge

Thank You



Department of Computer Science
California State University, Northridge